

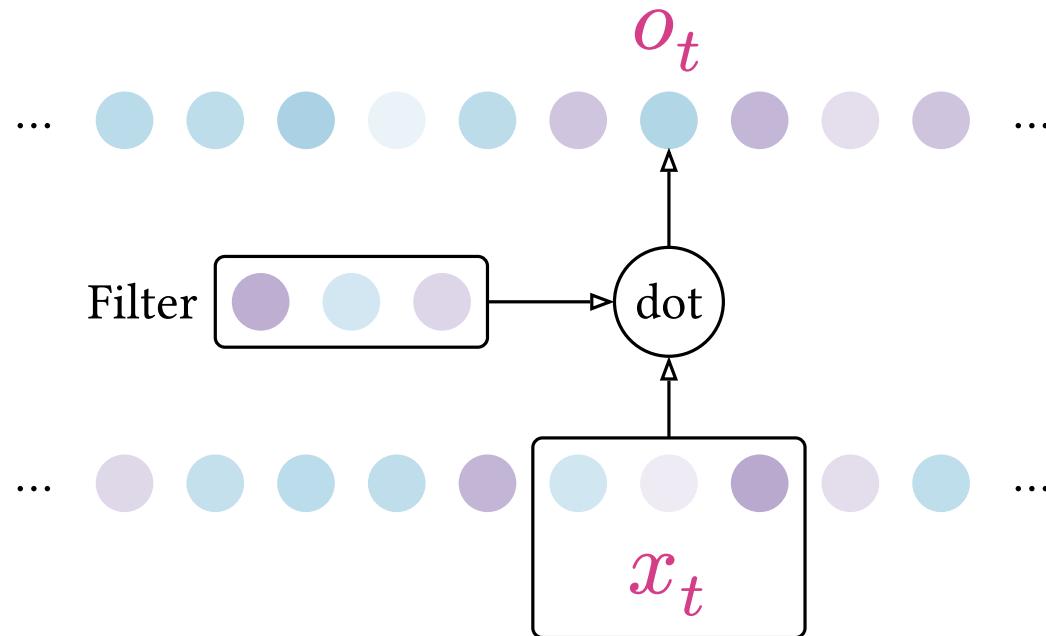
Deep Learning for Time Series

Session 2: ConvNets and Recurrent architectures

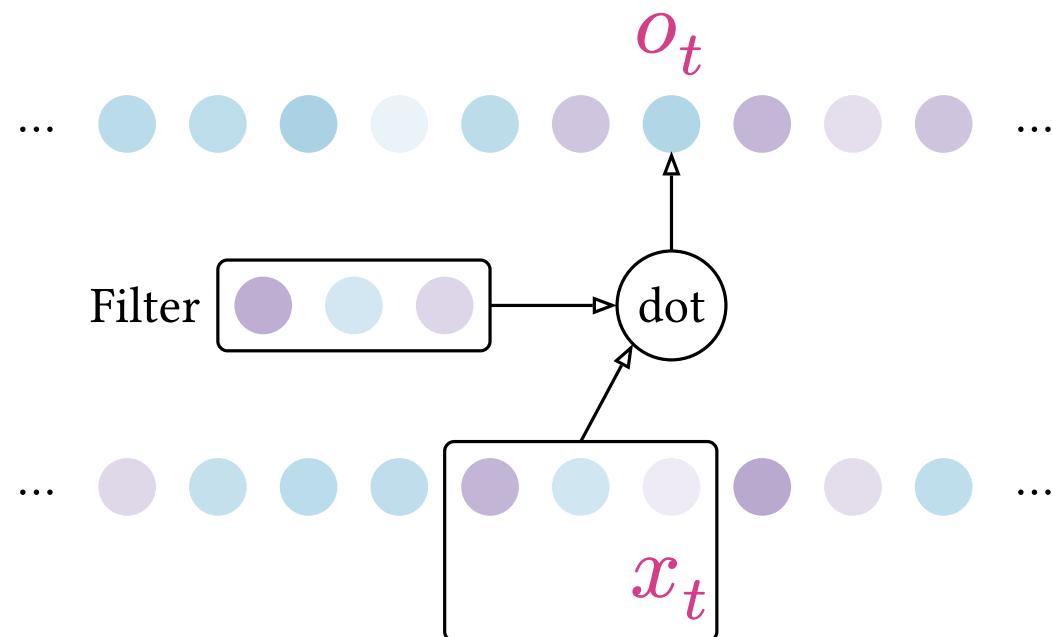
Romain Tavenard

Convolutional architectures

- Basic time series processing: 1d convolutions (over time)
- Limited receptive field: co-localization matters



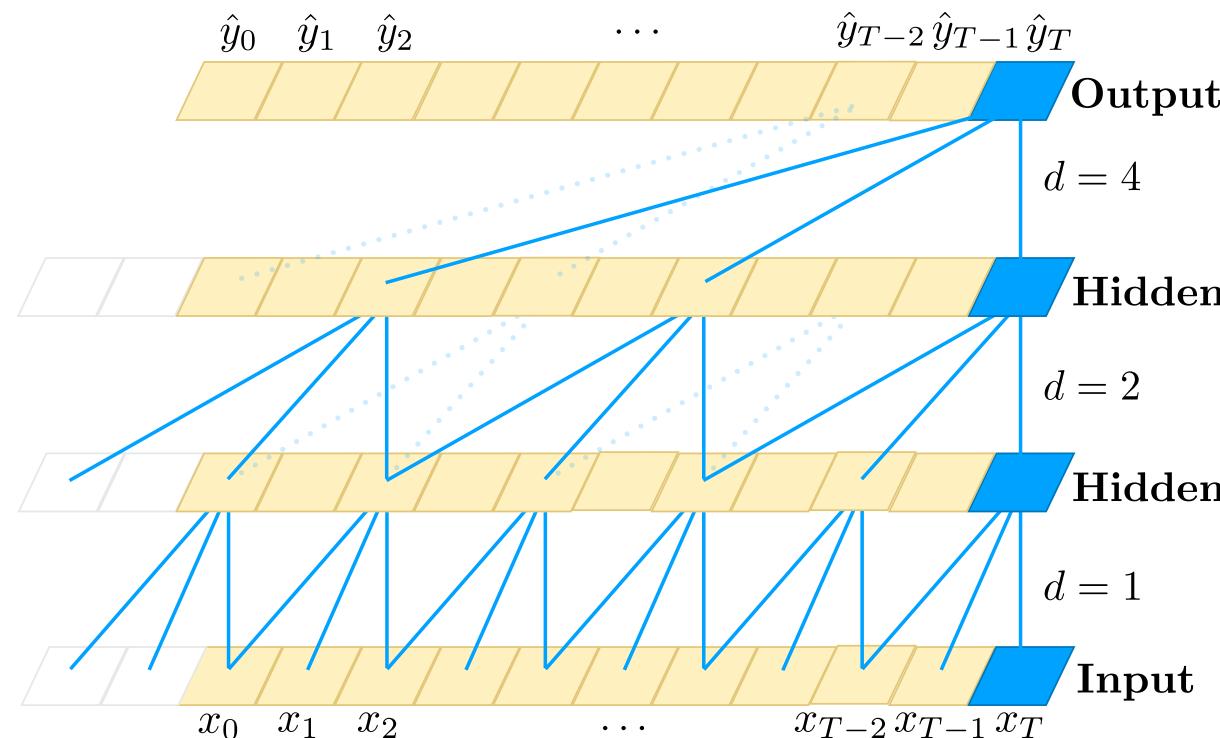
- Forecasting tasks: cannot access the future
- Causal convolution: convolve on past information alone (asymmetric window)



Temporal Convolution Network (TCN)

Convolutional architectures

- Main idea: cascade dilated causal convolutions
⇒ Larger receptive field

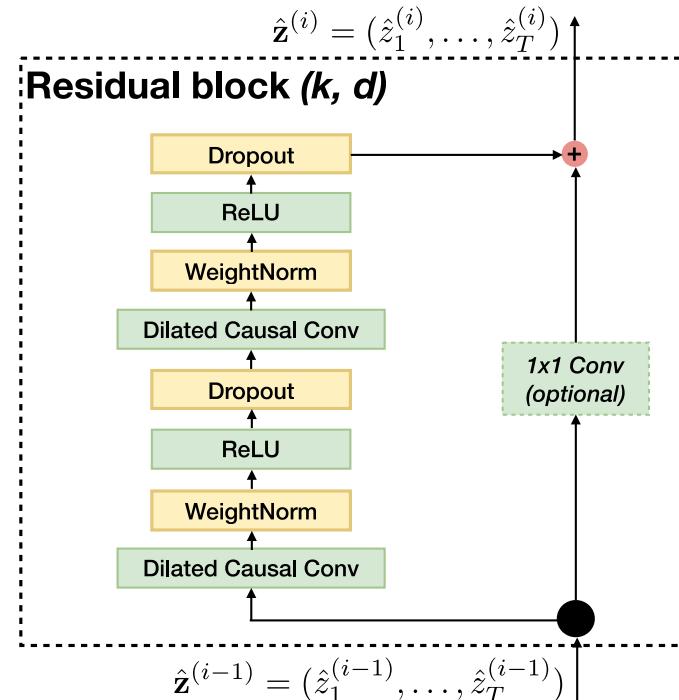


Source: “An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling”, Bai et al.

Temporal Convolution Network (TCN)

Convolutional architectures

- Additional improvements:
 - ▶ Residual connections
⇒ Multi-resolution analysis
 - ▶ Normalization+Dropout layers



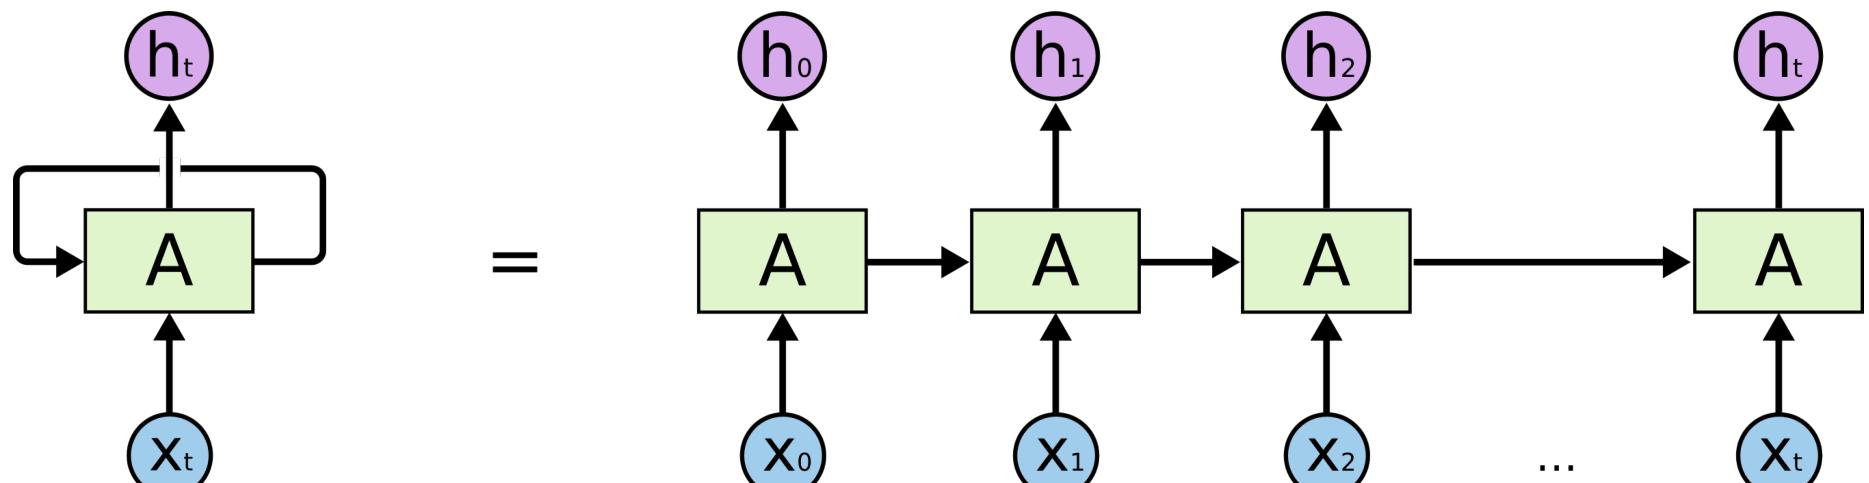
Source: “An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling”, Bai et al.

Recurrent architectures

Recurrent Neural Networks (RNNs)

Recurrent architectures

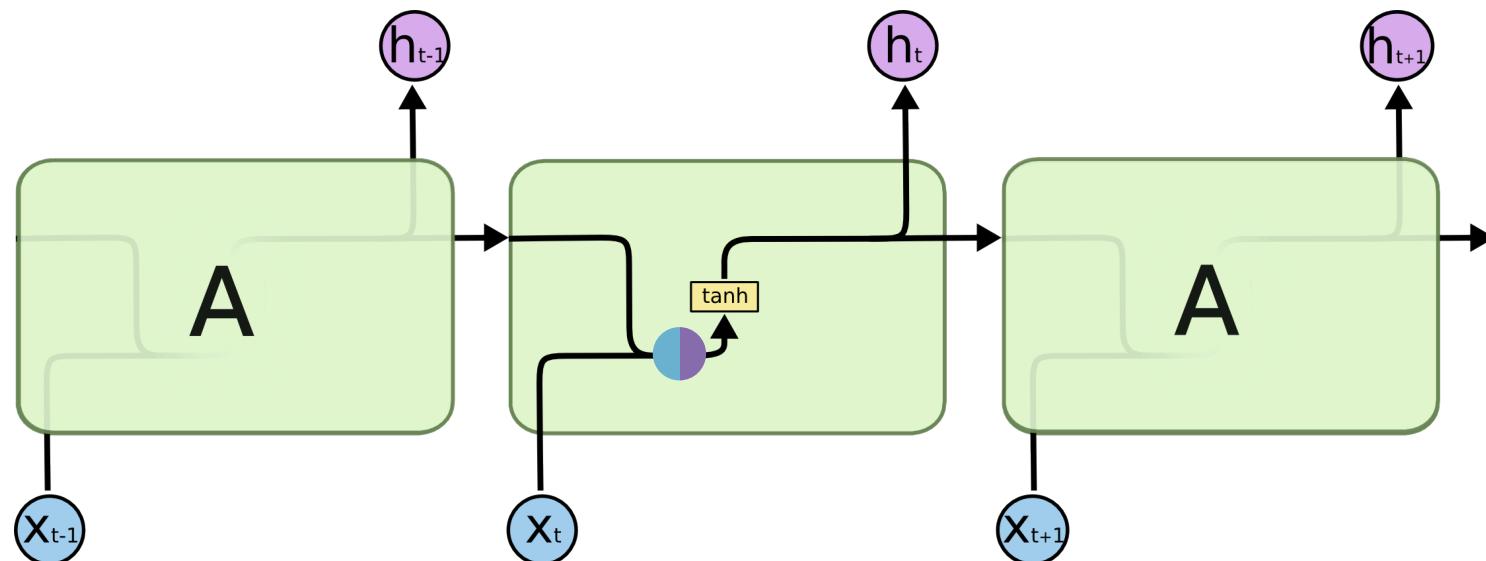
- Very flexible model (any length, let the model learn its memory needs, ...)



Source: Christopher Olah's blog

- Hidden state is computed as:

$$h_t = \varphi(x_t)$$

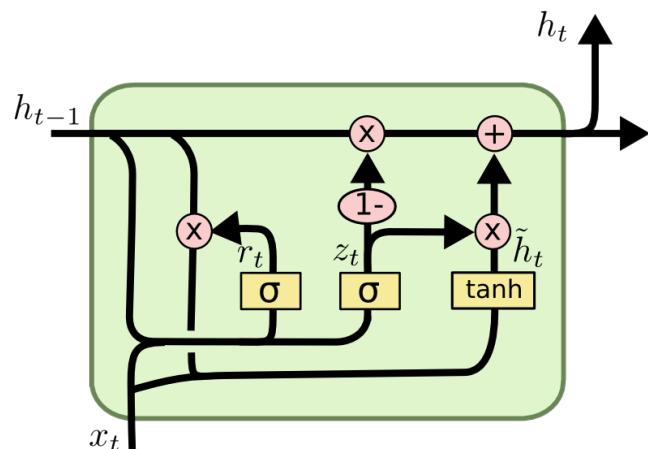


Source: Christopher Olah's blog

● x_t ● h_{t-1} ● Linear mixing of x_t and h_{t-1}

- Very flexible model (any length, let the model learn its memory needs, ...)
- Difficult to learn in practice
 - Slow (lack of parallelism)
 - Vanishing gradients (hard to learn long-term dependencies) or exploding gradients (if φ is unbounded)

- At each time step, keep only part of the information
 - ▶ Through **gating mechanism**



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

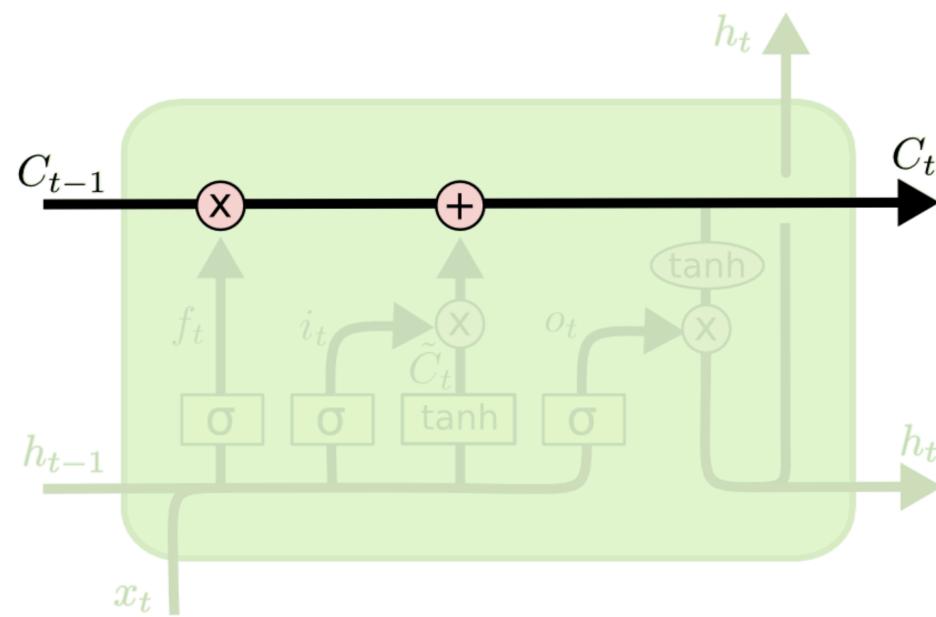
$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Source: Christopher Olah's blog

- Similar ideas as in GRUs, but:
 - ▶ an additional *cell state* C_t



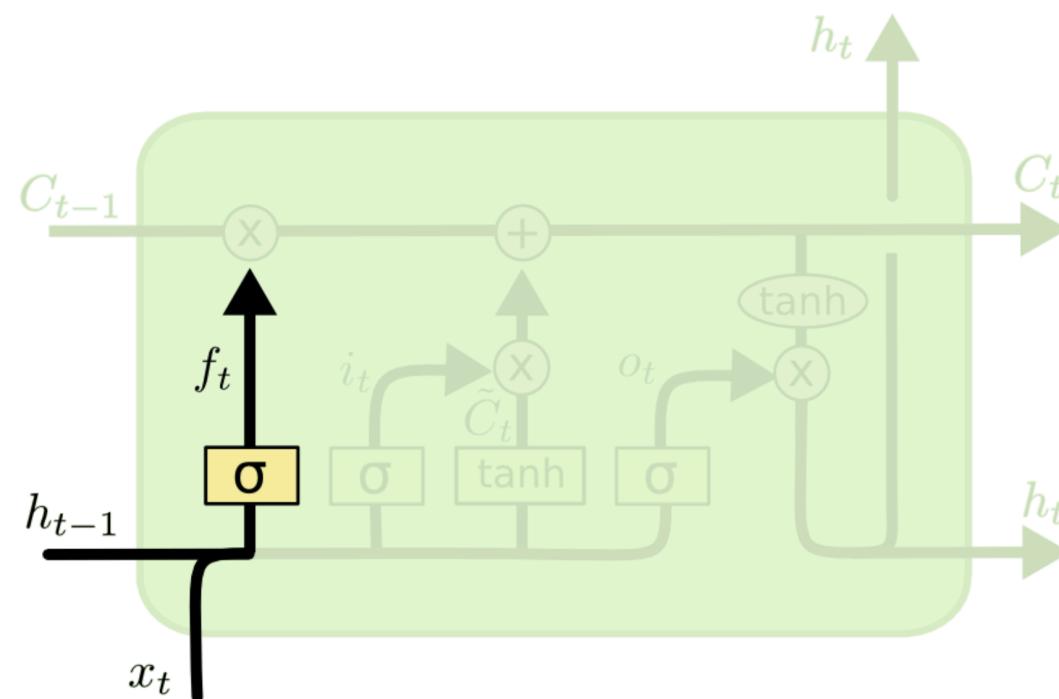
Source: Christopher Olah's blog

- ▶ input and forget gates are made independent
(in place of z_t in GRU)

Long Short Term Memory (LSTM)

Recurrent architectures

- **Forget gate:** $f_t = \sigma(\text{○})$



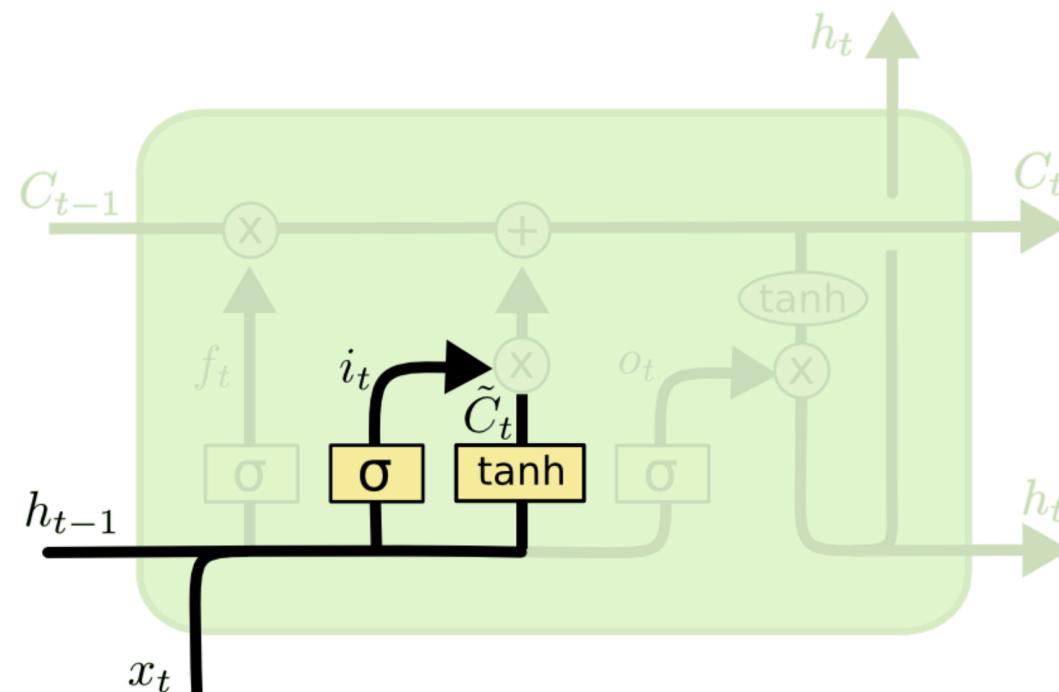
Source: Christopher Olah's blog

● x_t ● h_{t-1} ● Linear mixing of x_t and h_{t-1}

Long Short Term Memory (LSTM)

Recurrent architectures

- **Input gate:** $i_t = \sigma(\bullet)$
- **Suggested C_t update:** $\tilde{C}_t = \varphi(\bullet)$



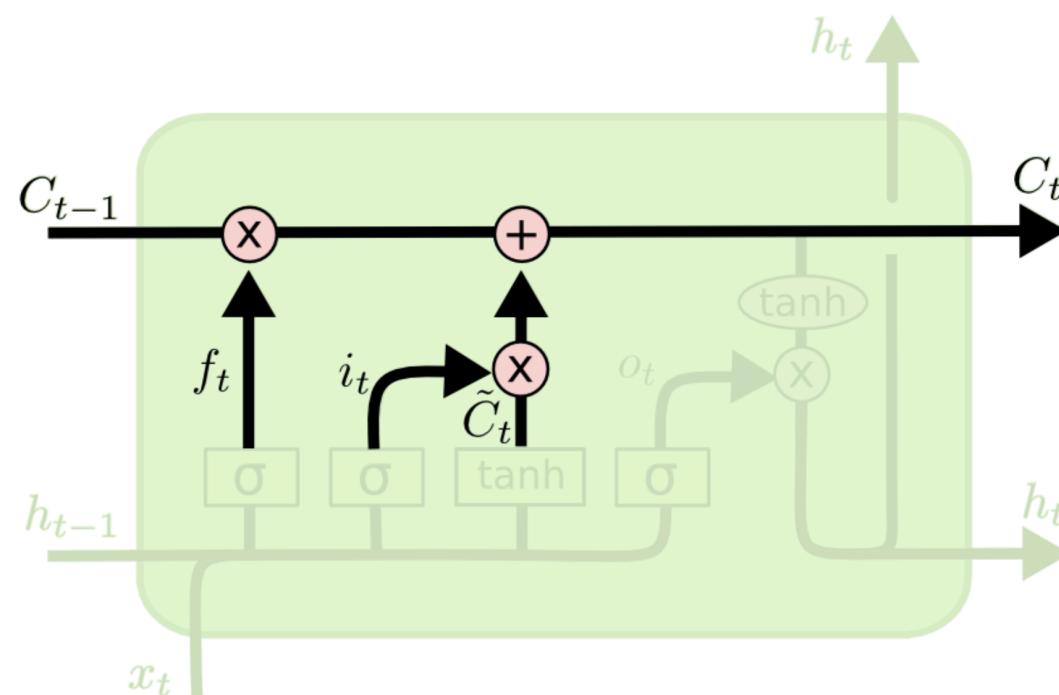
Source: Christopher Olah's blog

● x_t ● h_{t-1} ● Linear mixing of x_t and h_{t-1}

Long Short Term Memory (LSTM)

Recurrent architectures

- **C_t update rule:** $C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$

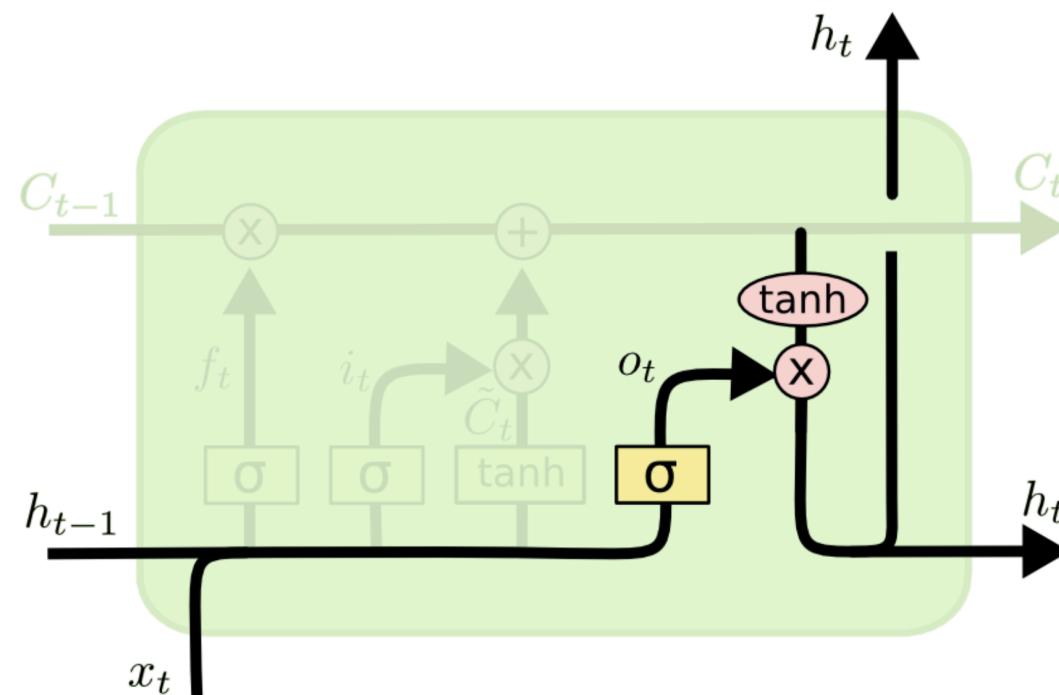


Source: Christopher Olah's blog

Long Short Term Memory (LSTM)

Recurrent architectures

- **Output gate:** $o_t = \sigma(\bullet)$
- **Hidden state update rule:** $h_t = o_t \odot \varphi(C_t)$



Source: Christopher Olah's blog



x_t

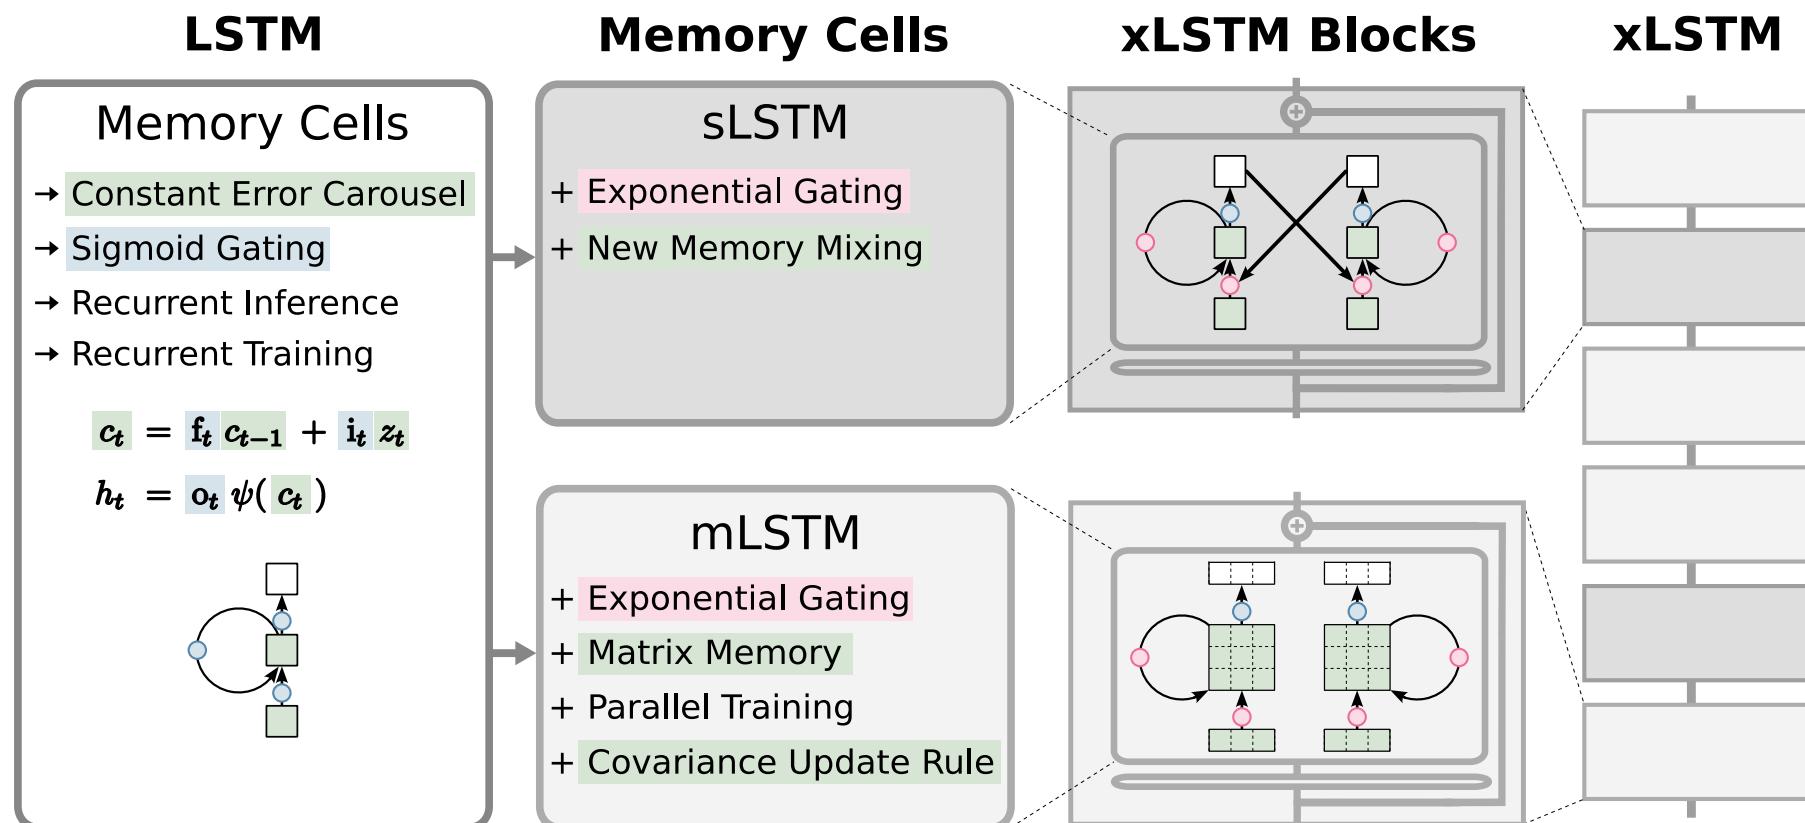


h_{t-1}



Linear mixing of x_t and h_{t-1}

- A “modern” LSTM variant
 - ▶ Made of sLSTM and mLSTM layers
 - ▶ Embedded in blocks with normalization layers, residual connections, *à la* Transformer



- What's “new”?
 - ▶ In both sLSTM and mLSTM layers:
 - Exponential activation (to face vanishing gradients)
 - ▶ In sLSTM only:
 - Multi-head
 - ▶ In mLSTM only:
 - Novel memory store
 - Drop recurrence for gate computations: better parallelism

- Exponential activation for input and forget gates:

$$i_t = \exp(\text{○})$$

$$f_t = \max(\exp(\text{○}), \sigma(\text{○}))$$

- Multi-head: keep separate linear mixings per head

● x_t ● h_{t-1} ● Linear mixing of x_t and h_{t-1}

- Exponential activation as in sLSTM
- Memory store

$$C_t = f_t \odot C_{t-1} + i_t \odot v_t k_t^\top$$

$$\tilde{h}_t = C_t q_t \quad (\text{up to normalization})$$

- ▶ Simplified case (no gate):
⇒ Mimicks QKV behaviour in self-attention
- Drop recurrence for gate computations: better parallelism

$$i_t = \exp(\bullet)$$

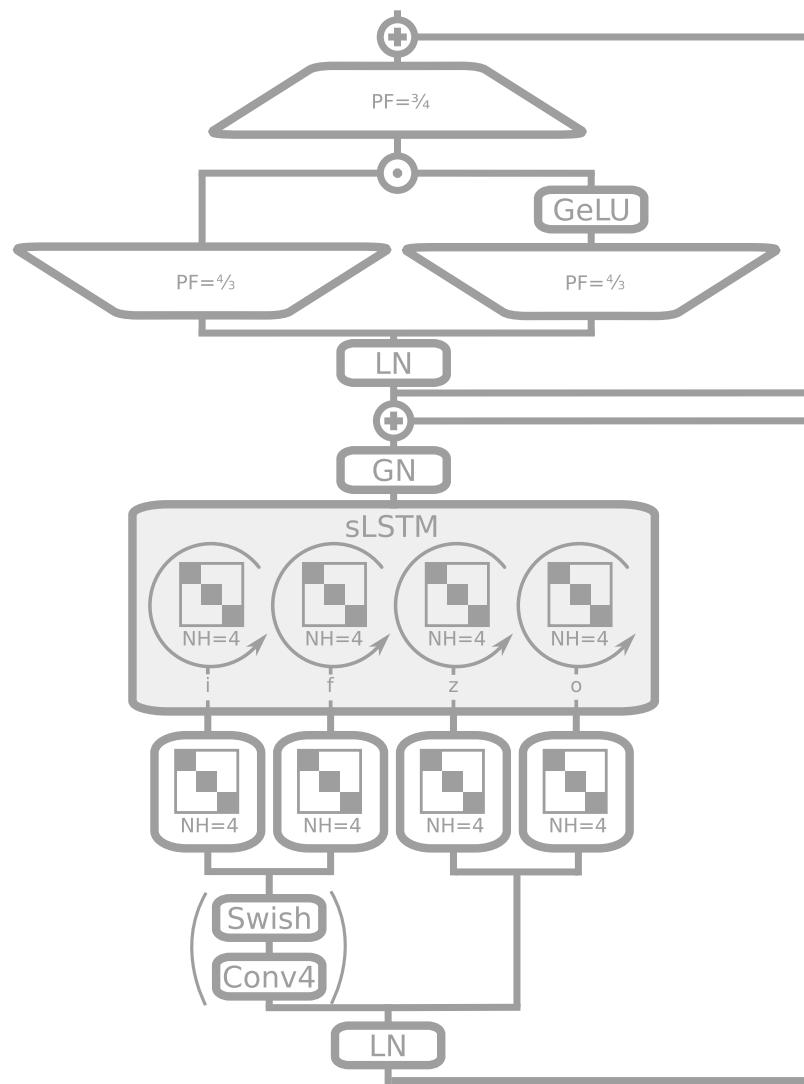
$$f_t = \max(\exp(\bullet), \sigma(\bullet))$$

$$o_t = \sigma(\bullet)$$

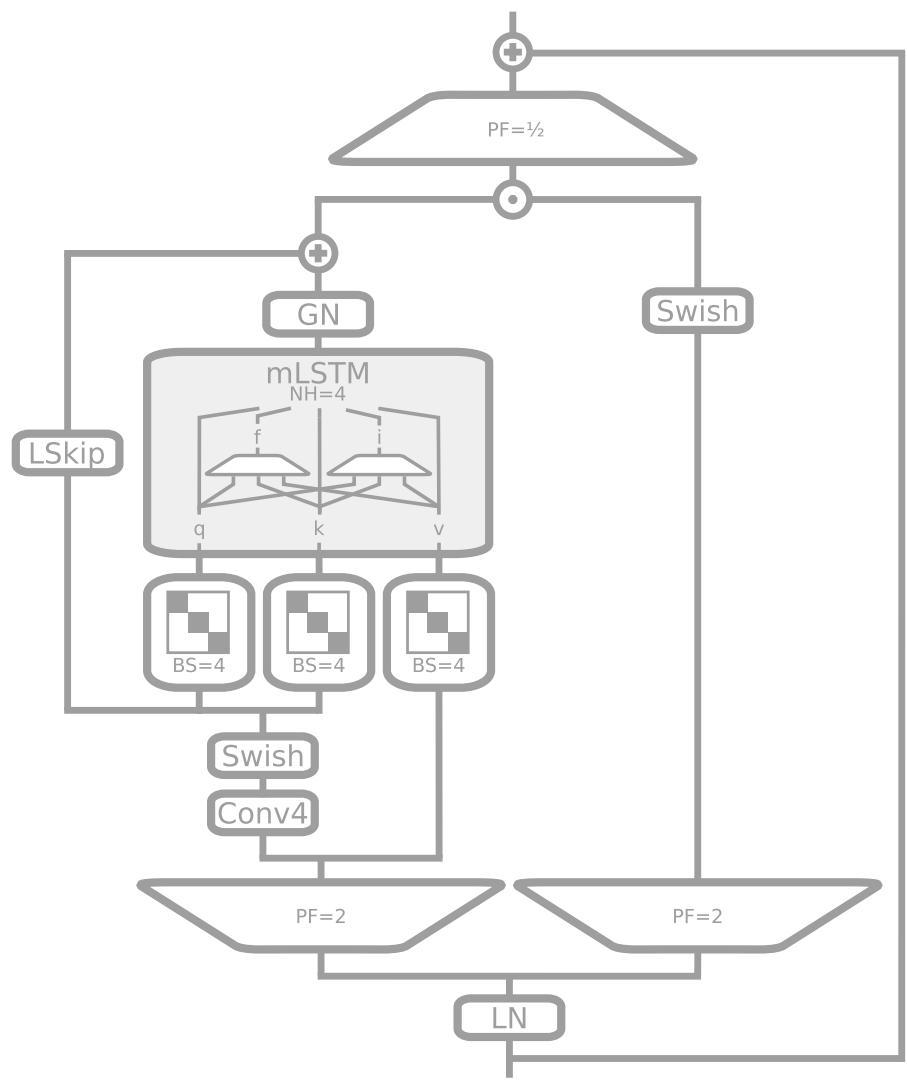
 x_t  h_{t-1}  Linear mixing of x_t and h_{t-1}

xLSTM: Building blocks

Recurrent architectures



An sLSTM block



An mLSTM block