# Deep Learning for Time Series

Session 4b: Time Series Foundation Models

Romain Tavenard
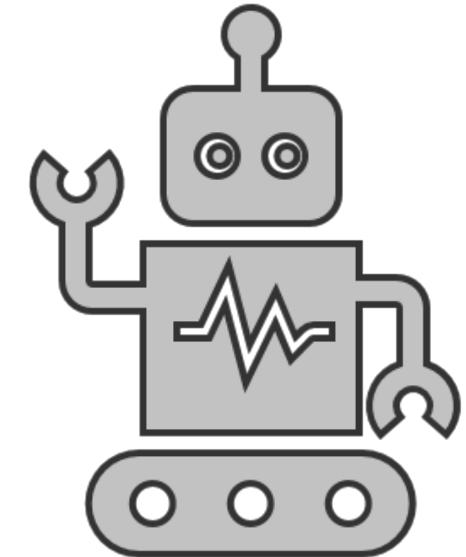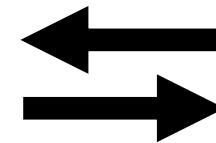
NB: Some figures in this slide deck are borrowed from Ievgen Redko's (great) course on time series forecasting and classification

# Traditional Machine Learning
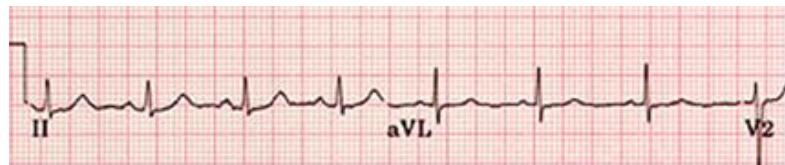
Normal ECG
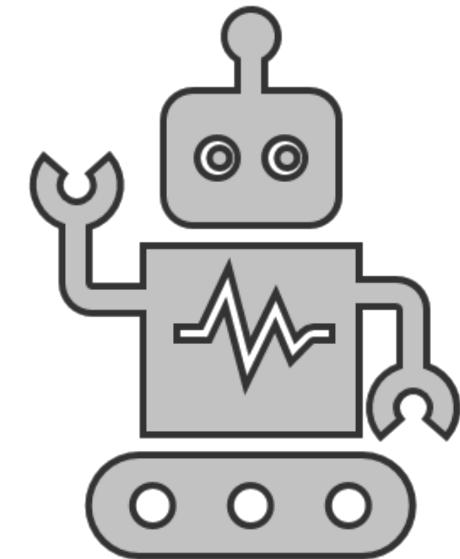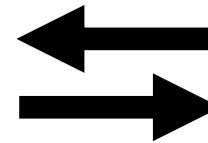


Abnormal ECG

- Traditional ML pipeline

  1. Collect training data

  2. Train a new model

  3. Deploy it

Traditional Machine Learning

Normal ECG


Abnormal ECG
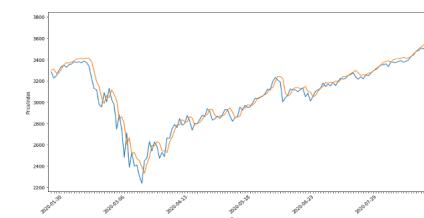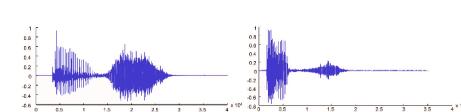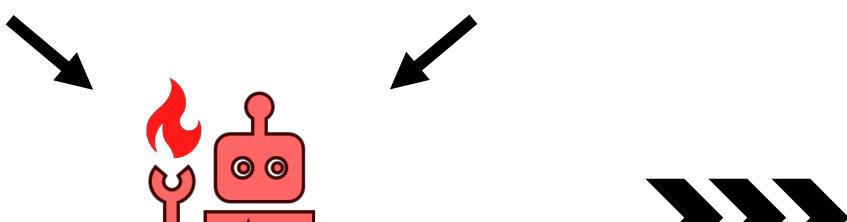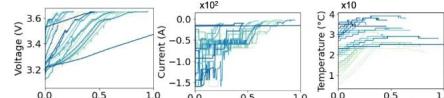
- Traditional ML pipeline

  1. Collect training data ← Issue 1: requires large training set

  2. Train a new model ← Issue 2: a new model for each task

  3. Deploy it

# Success of pre-training

- In Computer Vision: ImageNet pre-training is everywhere

- In text: multi-task LLMs

- In time series

  - Is there an ImageNet?

  - Do we have competitive task-adaptive models?

Time Series Foundation Models (TSFMs)
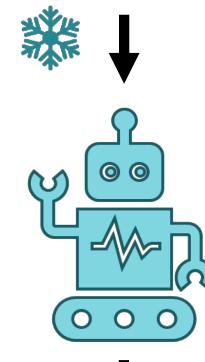
Step 1: Pre-training

Step 2: Use on New Tasks
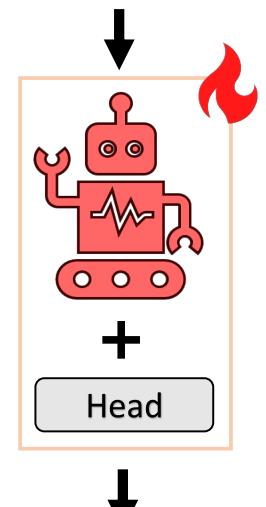
Option 1: zero-shot

Option 2: fine-tune

prediction

Head

prediction

Step 1: Pre-training

Pre-training datasets
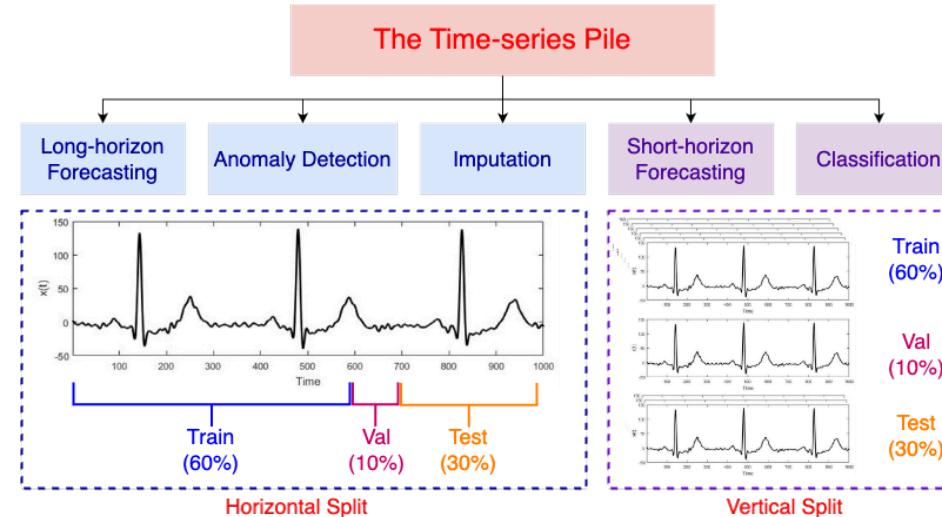
- **LOTSA (MOIRAI)**
  - 27.7B observations
  - 4M time series
  - LOTSAMini exists: < 1% of LOTSA
- **TimeSeries-PILE (MOMENT)**
  - 1.1B observations
  - 13M time series
- **Time300B**
  - 300B observations
  - 48M time series
- TimePFN's synthetic data generators
  - used 300-600M observations for TimePFN pre-training

# Pre-training strategies
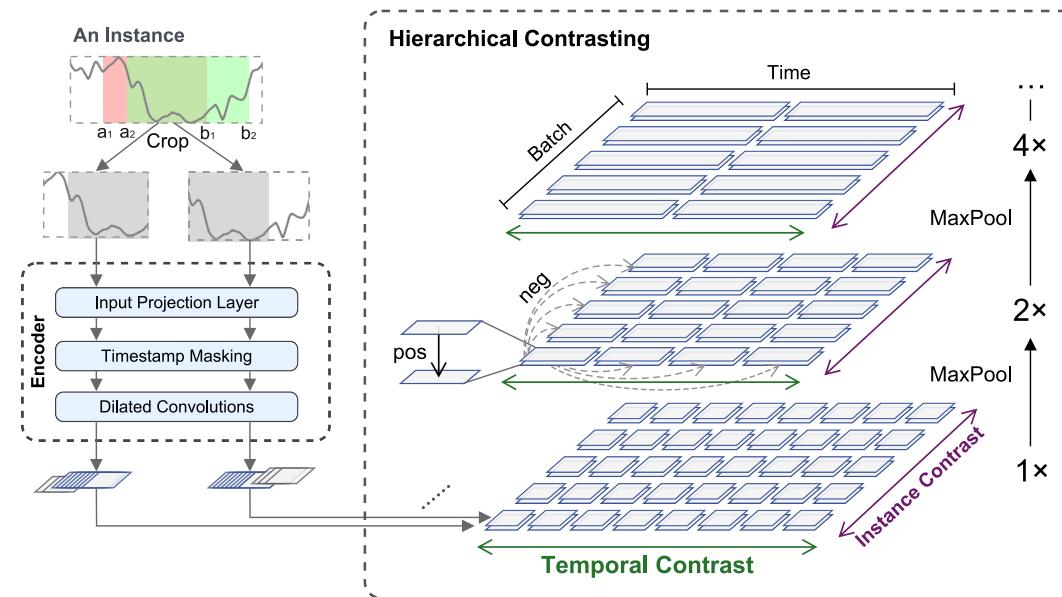
- Pre-training across tasks

Source: "MOMENT: A Family of Open Time-series Foundation Models",

ICML'24

- Semi-Supervised Learning
  - ‣ eg. MOMENT (masking), TS2Vec (contrastive)
- Use pre-trained LLM
  - ‣ eg. GPT4TS

- Self-supervised contrastive learning for time series
  - Contrast positives: same series, different segments/scales
  - Contrast negatives: different time series
- Learns multi-scale (hierarchical) representations
- Temporal encoder (CNN-based), no decoder
- Outputs fixed-length embeddings for downstream tasks
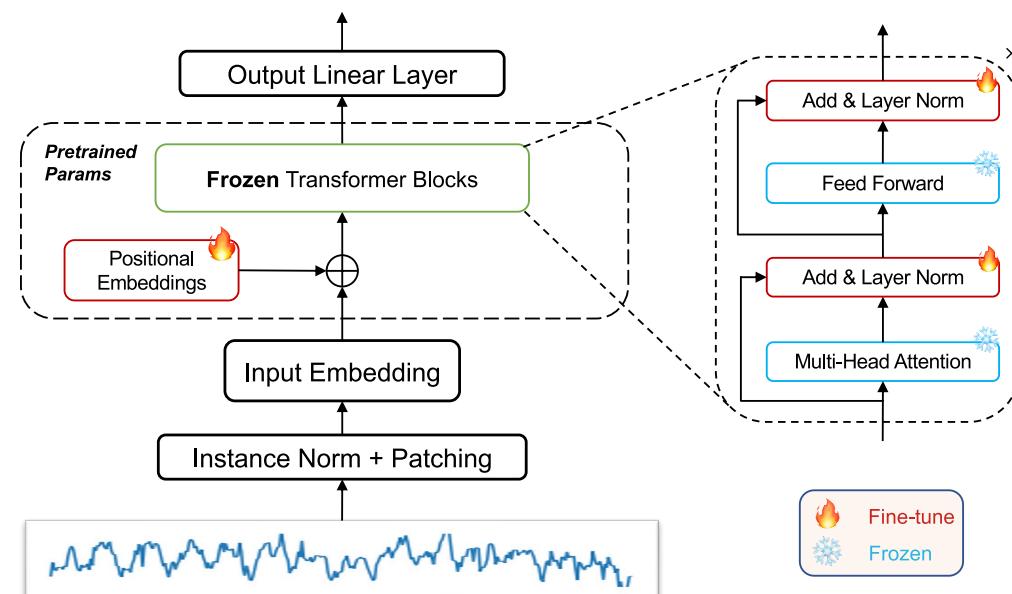
Source: "TS2Vec: Towards Universal Representation of Time Series", AAAI'22

- First strong task-agnostic TS representation method
- Validated pretrain once $\rightarrow$ reuse everywhere paradigm
- Shifted SSL from reconstruction to semantic contrastive learning
- Strong zero-shot / linear-probe baseline

General idea: make use of LLM backbone

1. Normalization + Patching (*à la* PatchTST)

2. Projection to LLM embedding dimension

3. Fine-tuning of positional embeddings + layer normalization, keeping others frozen

Source: "One Fits All: Power General Time Series Analysis by Pretrained LM", NeurIPS'23

Step 2: Adapting

MOIRAI
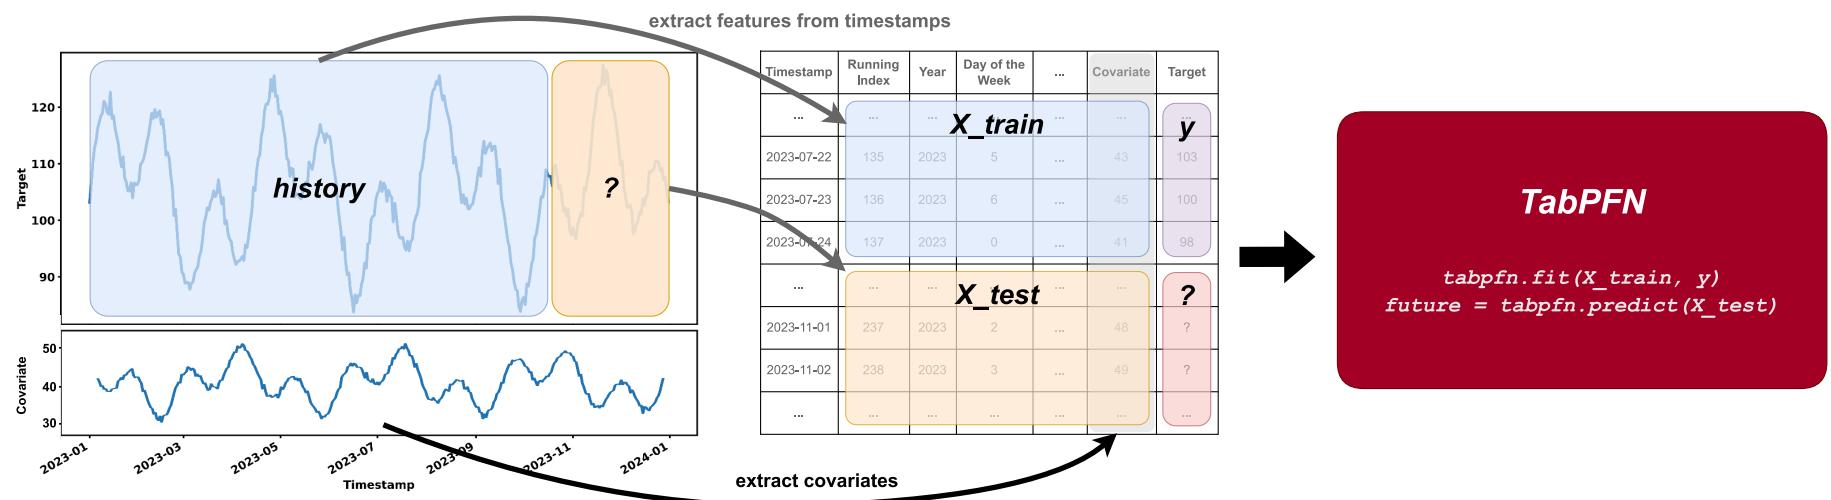
## Extract patches at different scales to catch multi-resolution information
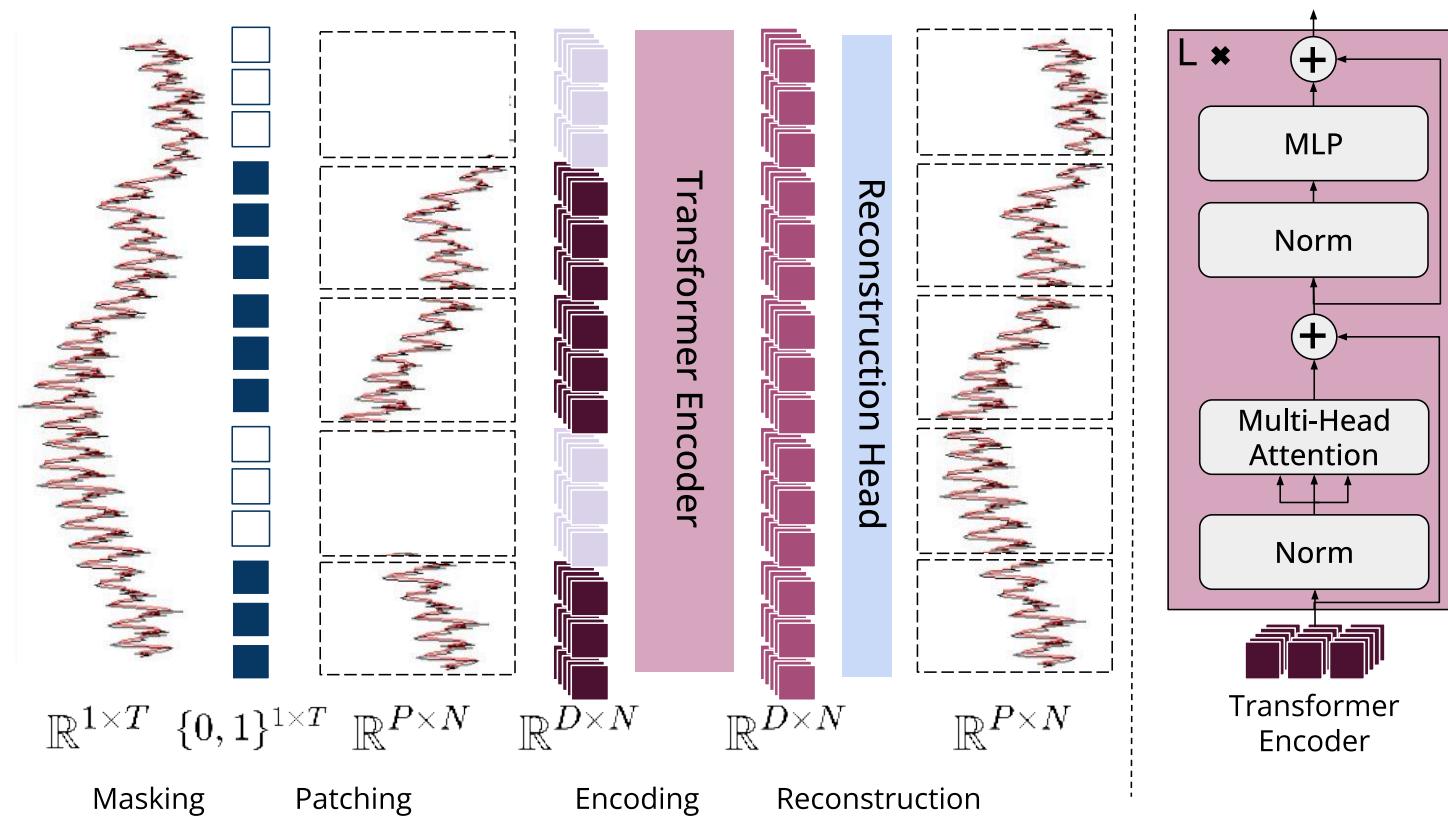
Source: "Unified Training of Universal Time Series Forecasting Transformers", ICML'24

- TabPFN is pre-trained on tabular data alone

- Builds basic TS features (calendar, seasonal, …)

- No TS prior at all!
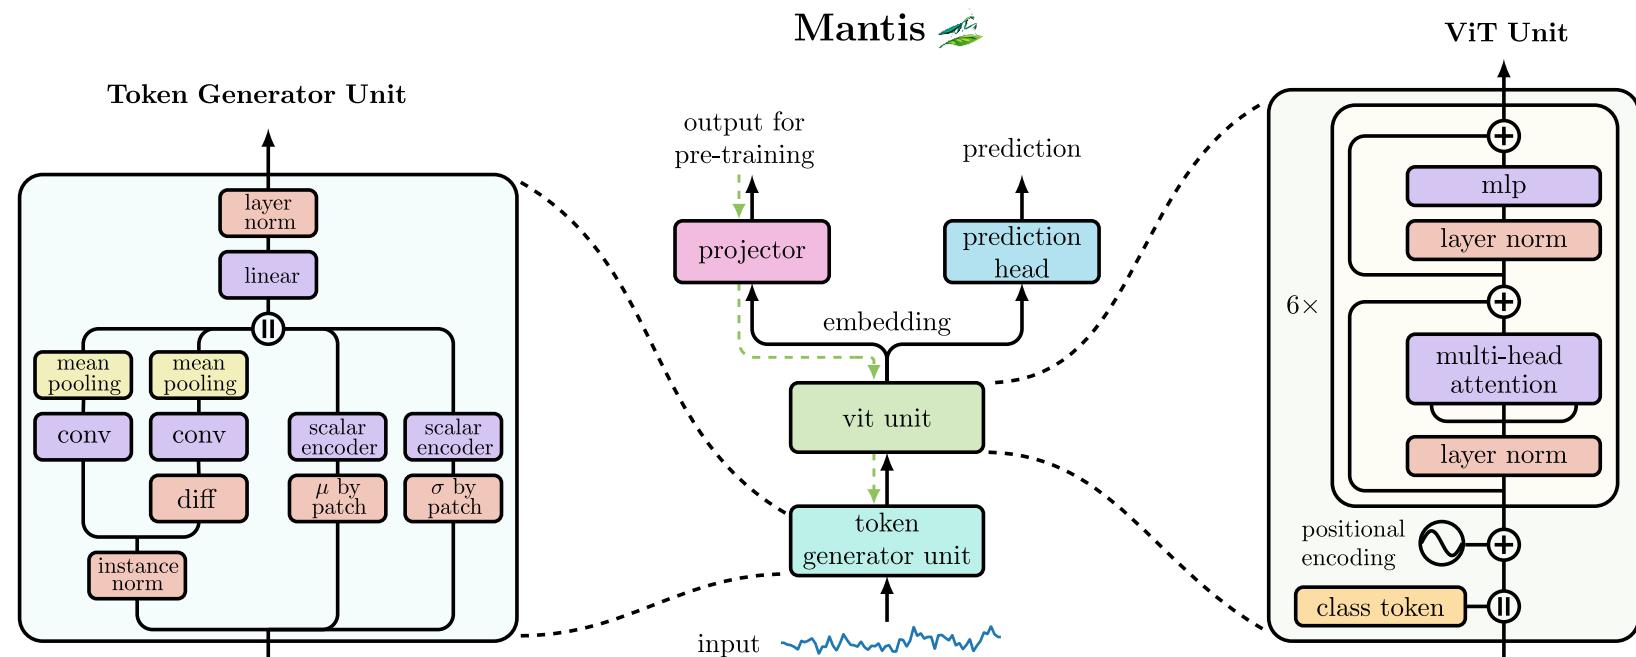
Source: "From Tables to Time: Extending TabPFN-v2 to Time Series Forecasting", ArXiV'26

Encoder-only model: can be used for any downstream task (classification, forecasting, imputation etc)
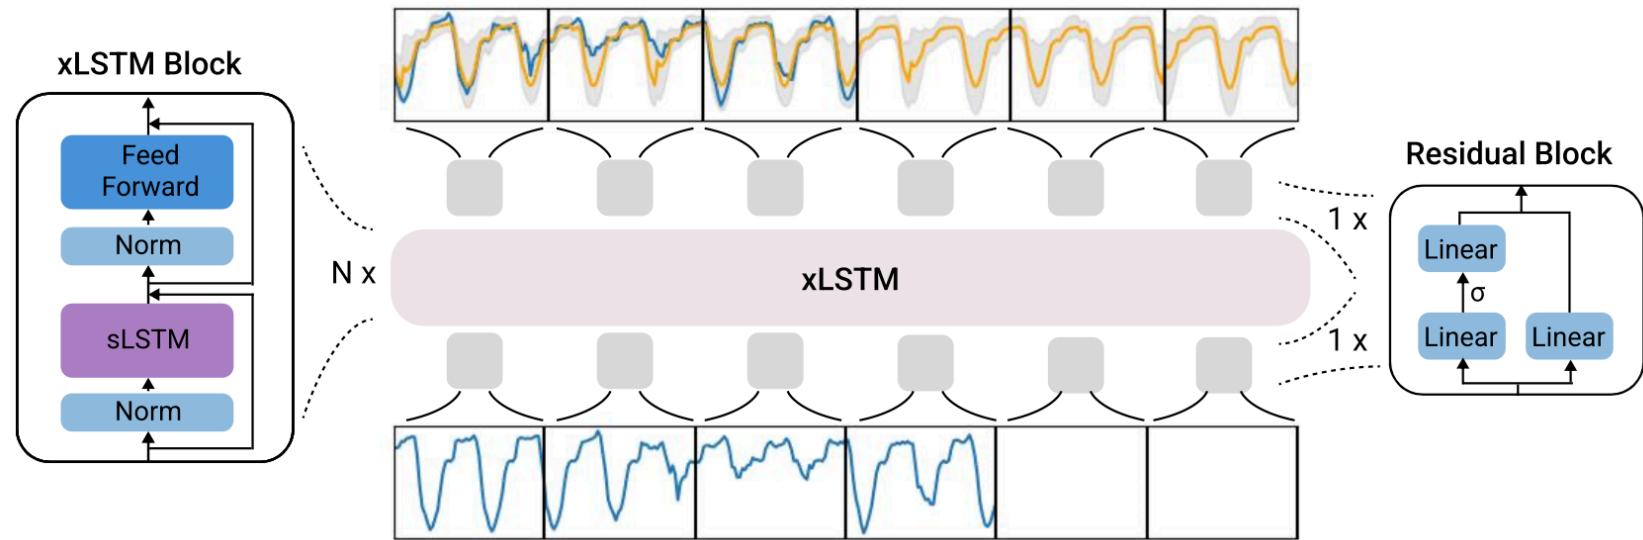

Source: "MOMENT: A Family of Open Time-series Foundation Models", ICML'24

- Contrastive pre-training

- Scalar embedding module for preserving basic statistics

- Strong zero-shot and fine-tuning performance for TSC

Source: "Mantis: Lightweight Calibrated Foundation Model for User-Friendly Time Series Classification", ICML workshops'25

- Masking-based pre-training
- Simple xLSTM architecture (only sLSTM blocks)

Source: "TiRex: Zero-Shot Forecasting Across Long and Short Horizons with Enhanced In-Context Learning", NeurIPS'25

Step 3: Evaluating

GIFTEval — Benchmarking Time-Series Foundation Models

- Unified evaluation benchmark for TSFMs

- Focuses on forecasting

- Diverse datasets across domains and time scales

  - ‣ short / long term

  - ‣ univariate / multivariate

- https://huggingface.co/spaces/Salesforce/GIFT-Eval